Learning Group-Disentangled Representation for Interpretable Thoracic Pathologic Prediction

1st Hao Li
Hohai University
Nanjing, China
lihao1998h@163.com

2nd Yirui Wu*
Hohai University
Nanjing, China
wuyirui@hhu.edu.cn

3rd Hexuan Hu
Hohai University
Nanjing, China
hexuan_hu@hhu.edu.cn

4nd Hu Lu
Jiangsu University
Zhenjiang, China
luhu@ujs.edu.cn

5nd Shaohua Wan
University of Electronic Science and Technology of China
Shenzhen, China
shaohua.wan@uestc.edu.cn

***Abstract*—Deep learning methods have shown significant performance in medical image analysis tasks. However, they generally act like "black box" without explanations in both feature extraction and decision processes, leading to lack of clinical insights and high risk assessments. To aid deep learning in envisioning diseases with visual clues, we propose Representation Group-Disentangling Network (RGD-Net), which can completely disentangle feature space of input X-ray images into several independent feature groups, each corresponding to a specific disease. Through complete disentanglement, RGD-Net improves both interpretability and accuracy in feature extraction and decision making for medical image analysis. Taking several semantically related and labeled X-ray images as input, RGD-Net firstly extracts completely group-disentangled representations of diseases through Group-Disentangle Module, which applies group-swap and linking operations to construct latent space by enforcing semantic consistency of attributes. To prevent learning degenerate representations defined as shortcut problem, we further introduce adversarial constricts on mapping from features to diseases, thus avoiding model collapse with former free-form disentanglement. Experiments on chestxray-14 and ChestXpert datasets demonstrate that RGD-Net are effective in predicting diseases with remarkable advantages, which leverage potential factors contributing to different diseases, thus enhancing interpretability in working patterns of deep learning methods.**

***Index Terms*—Interpretable Deep Learning, Group-Disentangled Representation Learning, Thoracic Pathologic Prediction, Adversarial Constricts**

I. Introduction

Despite deep learning methods have achieved remarkable progress in medical image analysis [1]–[3], most methods work as mappings from input factors to output classification results without explicit explanations. Therefore, they fail in promoting clinicians' and patients' confidence in trusting automatical diagnosis, thus preventing the usage of deep learning in medical domain.

Most attempts [4], [5] to explain deep learning focus on 'post-hoc' analysis by proving the importance of low-level visual features in producing accurate predictions. However, they couldn't directly link low-level visual features with high-level semantical diseases, and visually explain the decision making process. Essentially, both linking and explaining operations are valuable for clinicians to understand working patterns of deep learning for prediction.

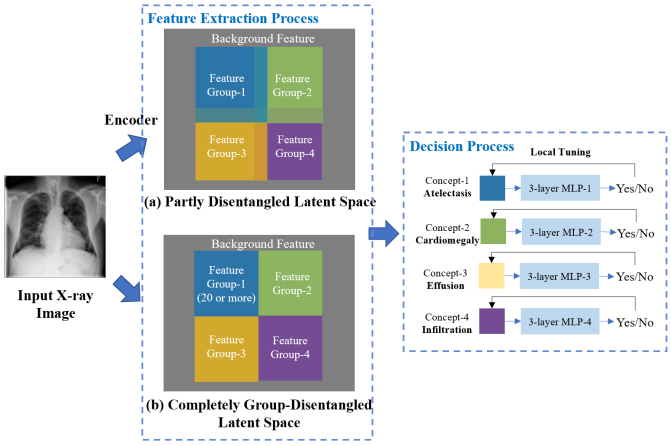

Fig. 1. (a)Partly disentangled latent space, where its feature groups are overlapped and coarse-grained, leading to confused explanations. (b)Completely disentangled latent space provided by RGD-Net, where feature groups are completely decomposed into independent subspaces, each of which corresponds to a specific disease.

As an alternative way, interpretable deep learning [6], [7] considers the inherent requirement of interpretation to embed clues based explanations in their neural network design. Most of them built their framework on variational auto-encoder (VAE), which achieve significant process towards explainable deep learning by performing linking and explaining steps with help of visual clues represented as feature groups. However, they generally ignore independence of learned clues, where they map visual samples onto a latent space that overlapped separates the information belonging to different attributes. Therefore, they only achieve partly disentangled effects with overlapping and coarse-grained low-level features as shown in Fig. 1(a), resulting in confused explanations and low accuracy classification results.

To achieve completely group-disentangled laten space as shown in Fig. 1(b), it's proved to enforce semantic consis-

tency of attributes, thus facilitating to leverage semantic links between samples. In other words, a completely disentangled latent representation space should be consist of subspaces, each encoding one attribute and each pair sharing no feature components. Following designs of completely disentanglement, training such a model usually faces one fundamental challenge, i.e., shortcut problem, that models may learn degenerate encodings by focusing on local minimum instead of global minimum, especially equipped with relatively free-form encoding network, such as VAE. Last but not least, how to develop task-specified interpretable deep learning methods remains an open question, due to the lack of involvement of existing clinical knowledge for either decision making or post explanation.

In this paper, we propose RGD-Net for interpretable thoracic pathologic prediction. We firstly achieve completely group-disentangled representations of diseases through the proposed Group-Disentangle Module. Such module is designed with group-swap and linking operations to leverage semantic links between input X-ray images and diseases, enforcing semantic consistency of attributes. To mitigate shortcut problem, we propose adversarial constricts, which borrows the idea of GAN to retain informative features during iteratively updating via group-swap and linking operations. Such constricts guarantee the model to seek for global minimum by forcing nash equilibrium between free-form grouping and convinced diagnosis, thus preventing model collapse. Significantly, we build a local-tuning medical application to showcase the power of interpretable thoracic pathologic prediction with RGD-Net, which could make reasonable decisions with part of subspace updating, relieving the computation burden of training from scratch when facing new samples or unsatisfactory results.

To sum up, our contributions are as follows:

- We propose *Representation Group-Disentangling Network* (RGD-Net), which completely extracts group-disentangled disease representations with fine-grained and non-overlapping features, thus promoting both interpretability and prediction accuracy.
- To resist shortcut problem caused by trapping in local minimum, an adversarial constraint is proposed to retain informative features during iteratively updating, thus forcing global minimum and avoiding model collapse.
- We experimentally demonstrate that RGD-Net can significantly improve classification accuracy, and showcase the potential local tuning medical application of RGD-Net, which not only enhances interpretable capability, but also relieves the burden of re-training.

II. Related work

The existing methods related to RGD-Net can be categorized into the following two types: Disentangled Representation Learning, and Thoracic Pathologic Prediction.

A. Disentangled Representation Learning

Disentangled Representation Learning [7]–[9] aims to separate the latent space of data into several parts, each representing a concrete, independent, and human-understandable concept. Existing methods for disentangling the latent space can be categorized into two classes: unit-disentanglement methods and group-disentanglement methods. The former methods treat a single latent unit as an independent concept. Following Variational Autoencoders [10], most unit-disentanglement methods [11]–[13] incorporate KL-divergence into the objective by forcing the latent factors to be statistically independent.

On the contrary, group-disentanglement methods treat a group of latent features as a concept. For example, Bouchacourt et al. [14] propose a deep probabilistic model to learn a disentangled representation from a set of grouped samples, and separate the latent representation into two swappable and semantical parts. In a similar way, Attila et al. [9] train their model with input triplets, in which the first two images should differ in the varying factor, but have the same common factor, and the third image is completely different from the first two. Furthermore, they offer solutions and analysis for shortcut problem and reference ambiguity in the disentangled representation learning.

Most relevant to our work, Group Supervised Learning (GSL) [15] is trained on groups of semantically related images and reconstruction objectives, allowing to decompose inputs into swappable components. Components from different images thus can be recombined to synthesize new samples. Unlike GSL, RGD-Net prefers implicit disease concepts rather than explicit attributes. Moreover, RGD-Net brings capability to be built in a practical medical application by solving the shortcut problem with adversarial constraints.

B. Thoracic Pathologic Prediction

To present ideas for predicting thoracic disease with latest improvement, we focus on related deep learning methods for readers' convenience. Early, Wang et al. [16] propose a weakly supervised framework for multi-label classification of chest diseases, which have done experiments on X-Ray8 dataset for 8 common chest pathologies. Then, Zhou et al. [17] propose a weakly supervised adaptive network, named as DENsenET-169, for chest disease recognition and classification in chest radiography. Specifically, they use different deep learning models for anomaly discovery classification and localization. Subsequently, Li et al. [1] propose a unified method for disease identification and localization with limited labeling data, where they adopt a Multi Instance learning (MIL) formula that improves performance compared with baseline models of ResNet and DenseNet.

Afterwards, Rajpurkar et al. [18] propose a 121-layer convolutional neural network CheXNet, which is trained on "ChestX-ray14", a expanded dataset of "ChestX-ray8" and containing over 100,000 frontal-view X-ray images with 14 diseases. Then, Wong et al. [19] propose a deep learning-based framework using Inception-ResNet-V2 for abnormal classification of chest X-ray images. Similarly, Wang et al. [20] propose a ChestNet model, which consists of a classification module and an attention module for computer-aided diagnosis of thoracic disease on CXR images.

Recently, one of the largest chest X-ray dataset "CheXpert" [21] is released, which contains 224,316 chest radiographs sampled from 65,240 patients. CheXpert provides researchers with an ideal volume of data for developing better and more robust supervised deep learning algorithms. Therefore, Irvin et al. [21] train a 121-layer DenseNet on "CheXpert" with various approaches to handle the uncertainty labels, where they finally achieve expert-level performance in detecting pathologies. Similarly, Hieu et al. [22] propose a multi-label classification framework based on deep convolutional neural networks, which diagnoses the presence of 14 common thoracic diseases on CheXpert. They exploit dependencies among abnormality labels via set of CNNs and use label smoothing regularization (LSR) for a better handling of uncertain samples.

III. Problem Definition

To improve interpretability of deep learning, auto-encoders are often used to provide hidden space decomposition, which is capable to obtain quantity of latent vectors containing vast of information corresponds to input chest X-ray images. Through reconstruction training, auto-encoder essentially compresses the input into the hidden space through encoder, and then reconstructs the original image through decoder based on hidden space.

Formally, we define Auto-Encoder : $\mathcal{X} \rightarrow \mathcal{X}$ as a combination of an encoder $E : \mathcal{X} \rightarrow \mathcal{R}^d$; and a decoder $D : \mathcal{R}^d \rightarrow \mathcal{X}$, where d denotes the dimension of the latent space $Z = E(X) \in \mathcal{R}^d$. To enhance interpretability in feature extraction, we wish to divide latent space into several semantic-specific parts as shown in Fig. 2. We define such property with the following formal definition.

Definition (Group-Disentangled Latent Space). *A group-disentangled latent space refers to a space consisting of several consecutive, non-overlapping subspaces, each of which is responsible for one specific concept.*

Such definition can also be expressed in the view of row-vectors:

$$z^{(1)} = [g_1^{(1)}, g_2^{(1)}, ..., g_m^{(1)}, b^{(1)}], \tag{1}$$

where row-vector $z^{(1)}$ is the concatenation of m row-vectors $\{g_i^{(1)} \in \mathcal{R}^{d_i}\}_{i=1}^m$ and a background row-vector $b^{(1)} \in \mathcal{R}^b$. It's noted that $d = \sum_{i=1}^m d_i + b$, where $\{d_i\}_{i=1}^m$ and b are hyper-parameters, and g_i corresponds to the concept c_i.

Although auto-encoder network can compress the image into the latent space and reconstruct it in a proper way, researchers still care about internal structure of latent vectors, since general auto-encoder fails in generating separate independent representations of concepts. Without such independent representations, users are hard to be convinced by external and explicit links between concepts and predictions. Therefore, we propose RGD-Net, which can completely disentangle the latent space into different subspaces, each part corresponding to a pathological concept. In that way, RGD-Net provides a robust and reasonable explanation on relationship between feature groups and disease labels.

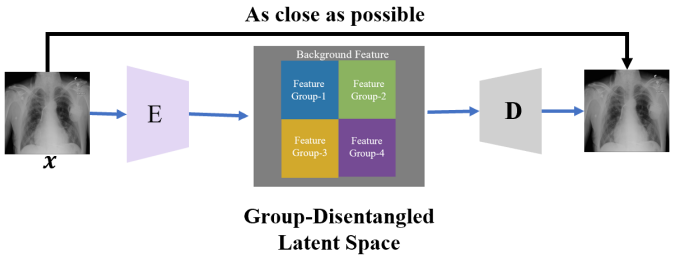

Fig. 2. Demonstration of encoding group-disentangled latent space based on auto-encoder. Each feature group in the latent space corresponds to a specific disease with medical usage.

IV. Methodology

We first introduce overall structure design of RGD-Net. Then, we describe in detail how we perform complete disentanglement via Group-Disentangled Module. Finally, we present the adversarial constrict that mitigates shortcut problem, even facing free-from swap and linking operations existed in former step.

A. Network Overview

As shown in Fig. 3, we propose RGD-Net to obtain group-disentangled latent space by completely disentangling representations of disease concepts (e.g., Atelectasis, Cardiomegaly, Effusion, and Infiltration in our case) based on a group of semantically related images. After training, we use the reconstructed group-disentangled latent space to perform quantity of downstream tasks, such as accurately predicting diseases based on testing images.

Specifically, RGD-Net firstly takes a group of semantically-related X-ray images as inputs. Then, it trains its encoder and decoder through the proposed Group-Disentangled Module and Adversarial Constrict, which forces the encoder structure to reconstruct the input image by a group-disentangled latent space. It's noted that Group-Disentangled Module contains two operations, i.e., Linking and Group-Swap Operation. Linking Operation acts like auto-encoder, which builds direct relationship between semantical concepts of disease and low-level visual features. We calculate its related reconstruction loss L_{lo} for each image. Meanwhile, Group-Swap Operation enforces semantic consistency of disease concepts constrained by the after-swap reconstruction loss L_{gs}. Last but not least, Adversarial Constrict builds on the idea of GAN by involving adversarial loss L_{ac} to solve the model collapse, that may encounter in the process of group-disentanglement and is generally defined as shortcut problem. During training, we combine three kinds of losses as a total loss L:

$$L = \min_{\mathrm{D,E}} \max_{Dis} L_{lo} + \lambda_{gs} L_{gs} + \lambda_{ac} L_{ac}, \tag{2}$$

where L_{lo}, L_{gs} and L_{ac} refer to losses of linking operation, group-swap operation and adversarial constrict part respectively, and scalar coefficients λ_{gs}, λ_{ac} represent the importance factor of different loss terms. It's noted that we optimize L by

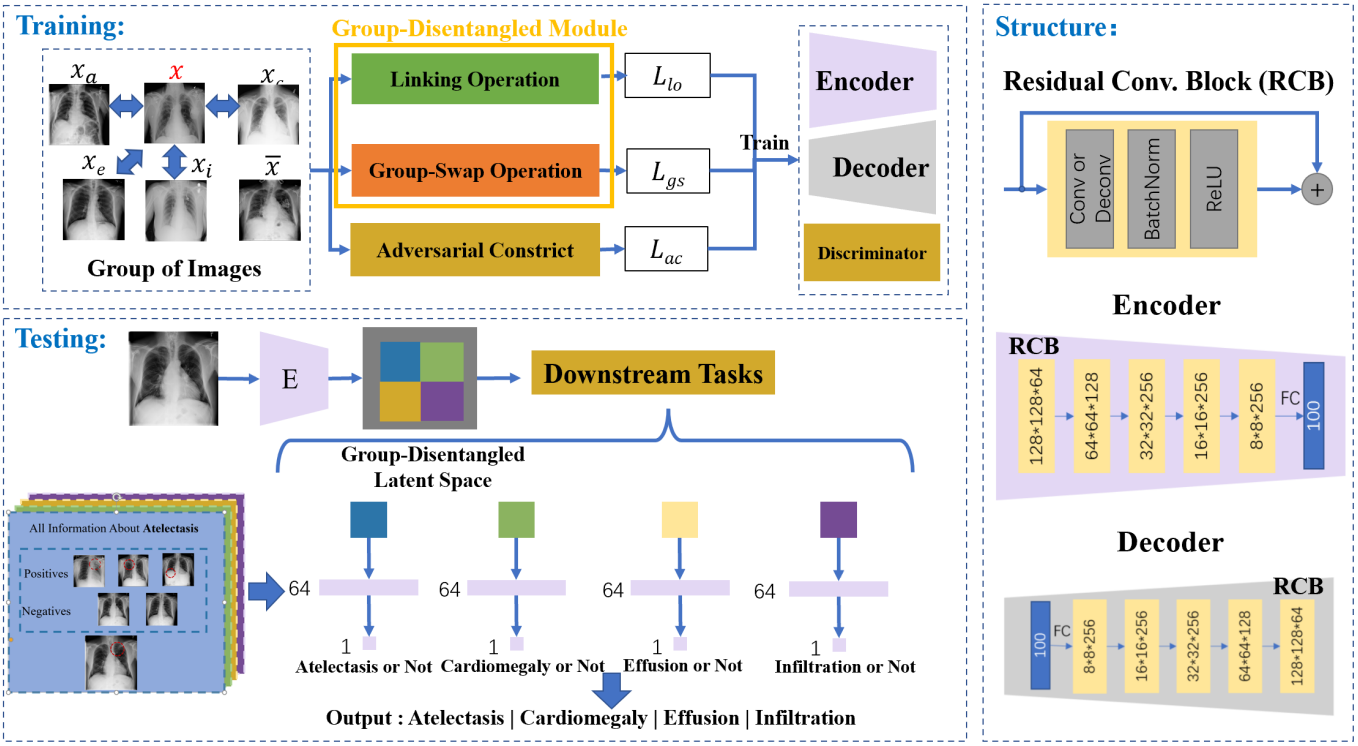

Fig. 3. The overall structure of RGD-Net, which extracts group-disentangled representations of disease through the Group-Disentangled Module and Adversarial Constrict. During testing, we use them to accurately predict corresponding disease labels.

gradient descent on parameters of encoder (E), decoder (D) and discriminator (Dis).

After training, it's supposed that we can apply RGD-Net on various computer vision tasks to provide convinced results. For example, it's expected to connect a classifier, e.g., MLP, with group of visual features for classification tasks. Alternatively, users can connect a detector with feature groups for detection tasks. In this paper, we demonstrate the effectiveness of RGD-Net to predict four categories of diseases based on chest X-ray images. Guided by the idea to apply on automatical medical application, we use a trained encoder to convert the input image into a group-disentangled latent space during testing phase. Afterwards, we predict thoracic pathologies disease concepts based on the new input X-ray images with an additional classification module with 3 layers of MLPs.

Encoder E is composed of a convolutional layer to generate feature map, four residual convolutional blocks with stride 2 for reshaping feature map to a vector, and a fully-connected layer to output a 100-dimensional vectors as latent feature. Meanwhile, decoder D mirrors the encoder in structure with a fully-connected layer, 4 residual de-conv blocks with stride 2 to reshape into a cuboid, and finally a de-conv layer to compute a synthesized image.

B. Group-Disentangled Module

While training our RGD-Net, we wish to group-disentangle these latent spaces by E. We use a group-swap operation based on auto-encoder to link low-level visual features with high-level disease concepts in the latent space.

As shown in Fig. 4 (a), we take a group of semantically related images S as the input of RGD-Net. First, we randomly select an image x from the data set, in this case, with atelectasis and cardiomegaly pathology. Based on the same pathology as x, we select images from the data set with Atelectasis and Cardiomegaly and without infiltration and effusion. Therefore, there are five images as input to RGD-Net. As input in this way, common properties between images can be effectively learned.

To retain the information of images in the hidden space, we use a auto-encoder based Linking Scheme, which links relationship between semantical concepts of disease and low-level visual features as shown in Fig. 4 (b). Specifically, for each input X, we embed data in a low-dimensional vector by the encoder. Then we link d_i units of the vector to a specific disease concept c_i. Formally, we select a subset of the latent space $g_i = [\mu_{l_i+1}, ...\mu_{l_i+d_i}]$, where l_i is the start position of the subset for concept c_i. Finally, we input this latent vector into the decoder and calculate the reconstruction loss L_{ls} for each image. As shown in Fig. 4 (c), we use the group-swap module to enforce semantic consistency of disease concepts, and extract features of disease concepts by leveraging semantic links between input images.

Taking an image pair sharing a disease as input, the group-swap module exchanges the corresponding part of the disease in the hidden space of the two images, and expects to get

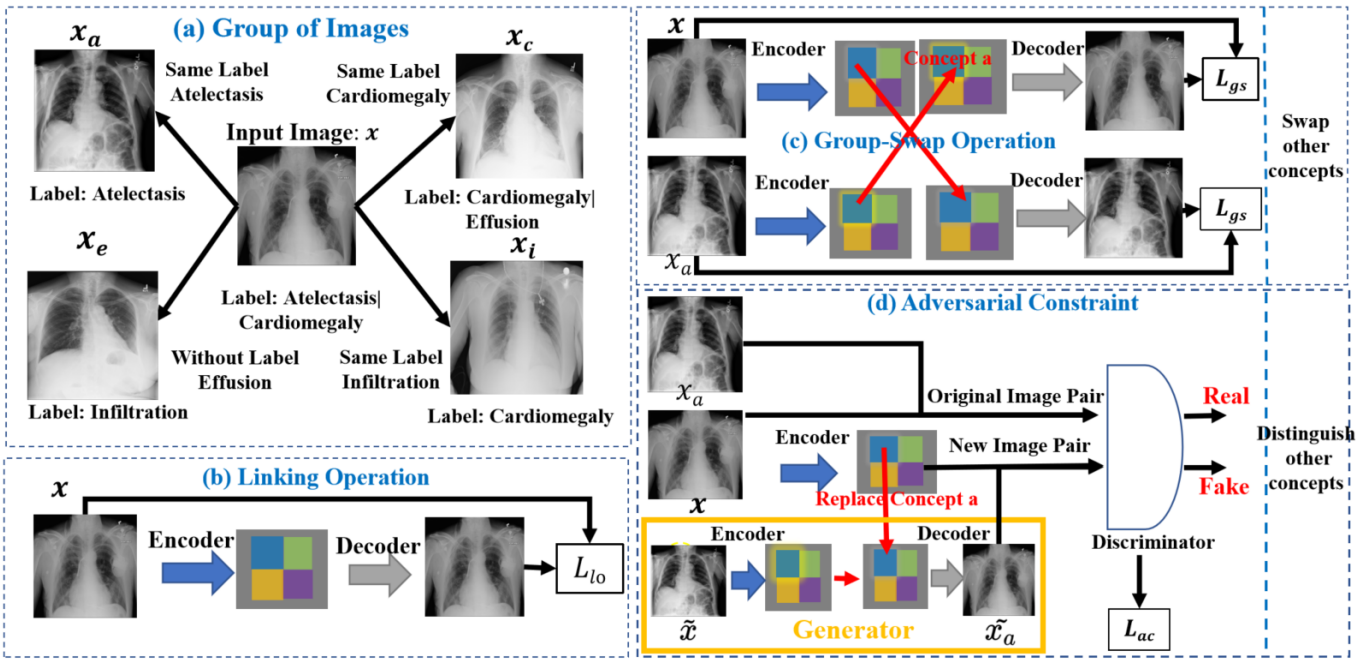

Fig. 4. Four steps in training RGD-Net to learn group-disentangled latent space: (a) **Group Images**, where we input a group of semantically related images to learn their common properties; (b) **Linking Operation**, where we link relationship between semantical concepts of disease and low-level visual features, calculating self-reconstruction loss for each image; (c) **Group-Swap Operation**: where we swap part of the latent representations of their shared concepts to enforce semantic consistency of disease concepts. (d) **Adversarial Constraint** takes triplet images as input to solve the *shortcut problem* caused by trapping in local minimal.

same result as the input through the decoder. Formally, for all $x^o \in S,\ x^o \neq x$, with the pair (x^o, x) share one concept value j (e.g., Atelectasis), the **Group-Swap** operation is defined as

$$z = E(x)\ ,\ z^o = E(x^o) \text{ and } \mathrm{z_s, z_s^o} = \textbf{swap}(\mathrm{z, z^o, j}), \quad (3)$$

where the **swap** operation is defined as

$$\begin{aligned} &\textbf{swap}(z^{(1)}, z^{(2)}, k) \\ &= \textbf{swap}([g_1^{(1)}, ..., g_k^{(1)}, ..., g_m^{(1)}, b^{(1)}], [g_1^{(2)}, ..., g_k^{(2)}, ..., g_m^{(2)}, b^{(2)}], k) \\ &= [g_1^{(1)}, ..., g_k^{(2)}, ..., g_m^{(1)}, b^{(1)}], [g_1^{(2)}, ..., g_k^{(1)}, ..., g_m^{(2)}, b^{(2)}]. \end{aligned} \quad (4)$$

The group-swap module is subject to the reconstruction loss

$$L_{gs} = ||D(z_s) - x||_2^2 + ||D(z_s^o) - x^o||_2^2. \quad (5)$$

C. Adversarial Constrict

Ideally, if there exist sufficient sample pairs sharing no duplicate concepts, loss of group-swap operation L_{gs} will be zero, so that complete group-disentanglement being logically obtained. However, due to free-form group-swap operation in former group-disentangled module, shortcut problem can occasionally occur with local minimum trap, where RGD-Net may learn degenerate encodings that all information of input images are retained in the group of background features.

It's supposed that shortcut problem can be mitigated by reducing dimensionality of background features, which not only forces the encoder to build a complete representation with several groups other than only one group of features, but also forces both encoder and decoder to properly link feature groups with the corresponding concepts under reasonable space room assumption. However, strategy of reducing dimensionality can only be convenient in practice, nevertheless resulting in time-consuming trial-and-error procedures to guess the proper dimensionality number.

Unlike adding constraints on feature space with setting a hyper-parameter, we propose an adversarial constrict to solve the shortcut problem without additional magic number to determine. As shown in Fig. 4 (d), we take triplet images, i.e., x_a, x and $\tilde{x_a}$, as input and introduce an adversarial training style. Specifically, the generator uses encoder-decoder structure to replace one specific feature group (represented as concept a) from x to $\tilde{x}$, thus generating new image $\tilde{x_a}$. In other words, the generator learns to generate an image containing pathological features shared by x and x_a, trying to fool the discriminator with new and fake image pair $[x, \tilde{x_a}]$. Meanwhile, the discriminator is designed as neural network to distinguish between original/real image pair $[x, x_a]$ and new/fake image pair $[x, \tilde{x_a}]$. Formally, $\tilde{x_a}$ can be defined as

$$\tilde{z} = E(\tilde{x})\ ,\ z = E(x) \text{ and } \tilde{\mathrm{x}}_\mathrm{a} = \textbf{swap}(\tilde{\mathrm{z}}, \mathrm{z}, \mathrm{a}). \quad (6)$$

In Fig. 4 (d), we show an example in adversarial training style by swapping the first feature group. Similarly, we construct image triples with different feature groups, i.e., concepts, and

calculate total adversarial losses with all image triples:

$$L_{ac} = \sum_{a=1,\ldots,5} log(Dis(x, x_a)) + log(1 - Dis(x, \bar{x_a})). \quad (7)$$

where the total number of disease concepts is 5 in our medical diagnosis application, and function $Dis()$ represents the discriminator to judge real or fake pair.

Essentially, the key idea in adversarial constricts is to achieve minimal adversarial loss, thus forcing informative low-level features belongs to disease concepts to keep remained during adversarial training process. In other words, global minimum can be achieved only if original image pair $[x, x_a]$ is real and new image pair $[x, \tilde{x_a}]$ is fake, so that replacing any concept can't beat the original concept pairing. In that way, we prove the best matching performance of the original image pair $[x, x_a]$, thus forbidding all image pairs collapse to wrong pairs. With global optimum of first adversarial loss L_{ac} and then total loss L, latent space could be completely and stably group-disentangled.

V. Experiments

We first evaluate the performance of our method on thoracic pathologic prediction and compare it with other non-disentangled DL methods. Then, to demonstrate the effectiveness of the module, we performed ablation experiments. Further, We evaluate our performance on learning group-disentangled representations and compare it with some partially disentangled and non-disentangled methods. Finally, we introduce the local tuning method to help clinicians improve performance in medical automated decision making.

A. Datasets and Measurements

We adopt two datasets to conduct thoracic pathologic prediction, i.e., chestxray-14 and ChestXpert. For the former dataset, we select a subset for experiments, which contains 36764 training images and 7353 testing images with 4 pathology labels (Atelectasis, Cardiomegaly, Effusion and Infiltration), which are extracting from the associated radiological reports using natural language processing. For the latter one, we also select a subset, which contains 162188 training images and 32437 testing images with 3 pathology labels(Pleural Effusion, Edema and Cardiomegaly).

To evaluate the performance of prediction, we follow the evaluation rules of both datasets, and adopt the area under receiver operating characteristic curve (AUROC) as our evaluation metric.

B. Accuracy of Thoracic Pathologic Prediction

Table. I shows that our RGD-Net, which has significantly improved on ChestXray-14 dataset by prediction with group-disentangled latent representation compared with the existing methods methods.

The AUROC values of the network on A, B and C reached 1, 2, and 3 respectively The AUROC values of RGD-Net on Atelectasis, Cardiomegaly, Effusion and Infiltration reached 86.30%, 89.80%, 92.69%, 86.53% respectively, being 5.36% , -2.68%, 6.31% and 13.08% higher than the second-highest achieved by CheXNet. Considering the reason for the decline of AUROC in predicting Cardiomegaly, we explored the ChestXray-14 and found that there was an extreme imbalance of the label of Cardiomegaly. This may be a weakness of interpretable models, making it difficult to learn concepts from these imbalanced datasets.

TABLE I
Comparison experiments on ChestXray-14 dataset. For each pathology, the highest AUROC scores are bolded.

Methods	**Atel**	**Card**	**Effu**	**Infi**
RGD-Net (ours)	**0.8630**	0.8980	**0.9269**	**0.8653**
CheXNet [18]	0.8094	**0.9248**	0.8638	0.7345
Yao et al. [23]	0.7720	0.9040	0.8590	0.6950
Wang et al. [16]	0.7160	0.8070	0.7840	0.6090
ChestNet [20]	0.7433	0.8748	0.8114	0.6772
Li et al. [1]	0.8000	0.8700	0.8700	0.7000
Zhou et al. [17]	0.8121	0.9066	0.8786	0.7065

TABLE II
Comparison experiments on ChestXpert dataset. For each pathology, the highest AUROC scores are bolded.

Methods	**Effu**	**Edema**	**Card**
RGD-Net (ours)	0.900	0.9023	0.8871
Ye et al. [24]	0.9166	0.9436	0.8703
Pham et al. [22]	**0.9640**	**0.9580**	**0.910**
Irvin et al. [21]	0.9360	0.9280	0.8540

To prove the performance of the proposed method on large medical datasets, we test the prediction performance of the proposed model on ChestXpert, one of the largest datasets currently available. As shown in Table. II, the accuracy of the proposed network is slightly lower on ChestXpert than the two latest networks, that is because our method considers not only the categories of predicted pathology, but also the interpretability of the network. It is useful in promoting clinicians' and patients' confidence and expanding the usage of DL in automated disease diagnosis. Moreover, our method can be easily migrated to other medical computing tasks. For example, we can replace the classification head with the detection head to detect the part related to the case.

C. Ablation Experiments

To verify the effectiveness of each module in the proposed method, we conduct ablation studies on ChestXray-14, where performance is listed in Table. III. The AUROC will decrease by a large percentage without the help of L_{gsm} module (row 3), since group-swap implies that swapping one attribute does not destroy latent information for other attributes. Moreover, the group-swap module can enforce semantic consistency of disease concepts, and extract features of disease concepts by leveraging semantic links between input images. When we remove this module, the model's understanding of the concept of disease is reduced, and it is naturally difficult to accurately predict the pathology.

TABLE IV

GROUP-DISENTANGLED REPRESENTATION ANALYSIS. WE USE THE ROW DISEASE FEATURES TO PREDICT THE AUROC OF COLUMN DISEASES ON THE TEST SET BY A SIMPLE 3-MLP. DIAGONALS ARE BOLDED AND '-' MEANS THAT ESTHER ET AL. [6] FAIL TO DISENTANGLE CONCEPT OF BACKGROUND.

	RGD-Net (ours) (completely disentangled)				Esther et al. [6] (partly disentangled)				AutoEncoder (without disentangled)			
Disease	Atel	Card	Effu	Infi	Atel	Card	Effu	Infi	Atel	Card	Effu	Infi
Atelectasis	**0.8630**	0.4855	0.5094	0.5005	**0.6136**	0.4960	0.4816	0.5050	**0.6076**	0.4990	0.4802	0.5297
Cardiomegaly	0.4822	**0.8980**	0.4836	0.5063	0.5062	**0.6610**	0.4968	0.4758	0.5067	**0.7048**	0.5183	0.4864
Effusion	0.4893	0.5061	**0.9269**	0.5229	0.5153	0.5038	**0.6688**	0.5099	0.4884	0.4985	**0.7444**	0.5292
Infiltration	0.4986	0.4900	0.4985	**0.8653**	0.4863	0.5230	0.5315	**0.5910**	0.4996	0.4955	0.4911	**0.6332**
Background	0.4983	0.5200	0.4926	0.4926	-	-	-	-	0.5045	0.5029	0.5087	0.4887

TABLE III

ABLATION AND DIVISION EXPERIMENTS. AUROC OF THORACIC PATHOLOGIC PREDICTION BY RGD-NET WITH DIFFERENT STRUCTURES ON CHESTXRAY-14.

Methods	Atel	Card	Effu	Infi
$L_{ls} + L_{gsm} + L_{GAN}$ (ours)	**0.8630**	0.8980	**0.9269**	**0.8653**
$L_{ls} + L_{gsm}$	0.6076	0.7048	0.7444	0.6332
$L_{ls} + L_{GAN}$	0.5263	0.5141	0.5365	0.5468
L_{ls} (Auto-Encoder)	0.5065	0.4713	0.5032	0.5289
Less Background	0.8497	0.8749	0.9013	0.8633
More Background	0.8263	**0.9048**	0.8883	0.8445

As removing the L_{GAN} module (shown in row 4), the model will degenerate to auto-encoder, the AUROC decreases again, indicating that the adversarial training module also has a certain group-disentanglement effect. As removing the L_{GAN} module from the RGD-Net (shown in row 2), AUROC drops to an approximate value of Auto-Encoder, which implies the fact that the model is collapsing. All the features of the image are kept in the background features, which makes the features not discriminative.

In the last two rows of Table. III, assuming that these thoracic pathologies are independent of each other, we distribute their corresponding latent subspace with same size. But these pathologies are not necessarily independent of each other and contain different amounts of information. Therefore, we modify the size of the latent space so that it is no longer equally divided. It's noted that less background of RGD-Net represents $g_i = 22, i = 1, 2, 3, 4$ and $b = 12$, and more background represents $g_i = 15, i = 1, 2, 3, 4$ and $b = 40$.

As we allocate less (as 12 in our paper) dimensions of latent space to represent background, the AUROC decreases by 1.33%, 2.31%, 2.56% and 0.2% for each pathology. If more (as 40 in our paper) latent space are used for background, the AUROC change by a percentage of -6.37%, +0.68%, -3.86% and -2.08%. This experiment shows that division equally is the most effective for this task.

D. Group-Disentangled Representation Analysis

To see the effect of group-disentanglement of our RGD-Net, we use the subspaces of disease concepts to predict four thoracic pathologies by a simple 3-MLP. If the hidden subspace contains all the information about the disease, the predicted result should be a matrix with 1 on the diagonal and 0.5 on the rest.

We use Esther et al. [6] and standard auto-encoder with classification head as comparison methods. The former partly disentangles the latent space, and the latter is not a disentangled method. Table. I shows that RGD-Net successfully decomposes the image into a group-disentangled latent space and uses each subspace to accurately predict the corresponding concept, but not to predict other concepts. Results of two comparison methods, whose latent space is not completely group-disentangled, show that each subspace doesn't know what it corresponds to, so their AUROCs are nearly 0.5.

This result shows that our method can effectively learn to group-disentangle representation and decompose the feature space into several independent parts, each of which represents a certain disease concept. However, other methods do not enforce the semantic consistency between the latent space and the concept of diseases, which leads to unsatisfactory results.

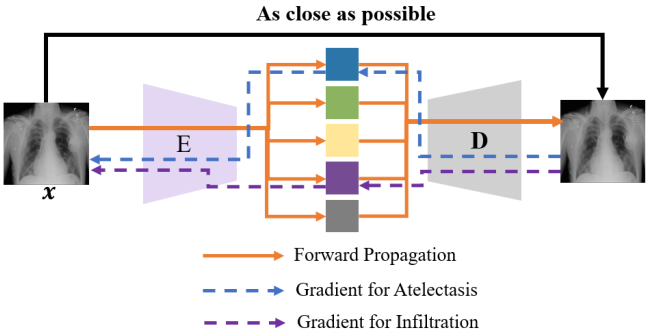

Fig. 5. Local tuning principle. When the model produces bad results in the medical automated decision-making process, the model can be tuned locally by limiting the gradient propagation, rather than tuning the whole model.

E. Local Tuning Medical Application

Based on RGD-Net, we can obtain a group-disentangled feature space, which not only brings interpretability and purer features, but also benefits medical automated decision-making processes. Consider yourself as a medical practitioner, while automating decisions with deep learning algorithms, has found that a few diseases are diagnosed with major errors, and that in most cases it works well. In this case, you should correct the part that caused the error and leave the rest unchanged.

Based on this idea, we propose a local tuning method for fine-tuning specific parts of the model in medical automated decision making, rather than retraining the entire model. As shown in Fig. 5, we use solid and dashed lines to represent the forward propagation and gradient back-propagation of the network respectively.

The forward propagation of the network is no different from other neural networks, but the specific diseases can be considered in the gradient back-propagation, and the gradient is calculated only for the corresponding part of the group-disentangled feature space, while the other parts are frozen. For example, when it is found that the model has a large error in the diagnosis of Atelectasis in the process of medical automated decision making, we can only make local tuning to the network, i.e., only allow the gradient of the corresponding features of Atelectasis to be transmitted, while gating the gradient of other features.

F. Implement Details

All our experiments were conducted on a server with two Intel Xeon E5-2620 v4 (@2.1GHz) CPUs and 4 NVIDIA GTX1080Ti graphic cards. Our experimental codes are mainly based on the PyTorch framework. Our initial learning rate is set as 0.0001, weight decay is 0.0001 and the momentum is 0.9. Due to the linear warm up mechanism, the learning rate increases from 1/30 to 0.01 in the first 500 iterations.

By default, both scalar coefficients λ_{gsm} and λ_{GAN} are set to 1. All the proposed modules are added, and the latent space is equally divided, e.g., if $d = 100$, then $d_i = 20, i = 1, 2, 3, 4$, and $b = 20$. For comparison purposes, the start positions for concepts are evenly distributed, i.e., $l_1 = 0$, $l_2 = 20$, $l_3 = 40$, and $l_4 = 60$.

VI. Conclusion

This paper proposes a Representation Group-Disentangling Network (RGD-Net), which completely extracts group-disentangled disease representations with fine-grained and non-overlapping features, thus promoting both interpretability and prediction accuracy. With clinical knowledge on two samples sharing the identical concept values, an group-swap module is introduced to link low-level visual features with high-level disease concepts in the latent space, thus laying an disease interpretable basis in feature extraction process. Further, we found the possible model collapse problem in the training process, and proposed an adversarial constrict to solve it. Finally, we experimentally demonstrate that RGD-Net can significantly improve classification accuracy compared with partly disentangled interpretable or other DL methods, and showcase the potential of RGD-Net to help clinicians improve performance in medical automated decision making by local tuning.

References

[1] Zhe Li, Chong Wang, Mei Han, Yuan Xue, Wei Wei, Li-Jia Li, and Li Fei-Fei, "Thoracic disease identification and localization with limited supervision," in *Proceedings of CVPR*, 2018, pp. 8290–8299.

[2] Leo K. Tam, Xiaosong Wang, Evrim Turkbey, Kevin Lu, Yuhong Wen, and Daguang Xu, "Weakly supervised one-stage vision and language disease detection using large scale pneumonia and pneumothorax studies," in *Proceedings of MICCAI*, 2020, vol. 12264, pp. 45–55.

[3] Faizan Karim, Munam Ali Shah, Hasan Ali Khattak, Zoobia Ameer, Umar Shoaib, Hafiz Tayyab Rauf, and Fadi Al-Turjman, "Towards an effective model for lung disease classification: Using dense capsule nets for early classification of lung diseases," *Appl. Soft Comput.*, vol. 124, pp. 109077, 2022.

[4] Bolei Zhou, Aditya Khosla, Àgata Lapedriza, Aude Oliva, and Antonio Torralba, "Learning deep features for discriminative localization," in *Proceedings of CVPR*, 2016, pp. 2921–2929.

[5] Ramprasaath R. Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra, "Grad-cam: Visual explanations from deep networks via gradient-based localization," in *Proceedings of ICCV*, 2017, pp. 618–626.

[6] Puyol-Antón Esther, Chen Chen, and James R. Clough, "Interpretable deep models for cardiac resynchronisation therapy response prediction," in *Proceedings of MICCAI*, 2020, vol. 12261, pp. 284–293.

[7] Matthew J. Vowels, Necati Cihan Camgöz, and Richard Bowden, "Gated variational autoencoders: Incorporating weak supervision to encourage disentanglement," in *Proceedings of 15th IEEE International Conference on Automatic Face and Gesture Recognition*, 2020, pp. 125–132.

[8] Yoshua Bengio, Aaron C. Courville, and Pascal Vincent, "Representation learning: A review and new perspectives," *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 35, no. 8, pp. 1798–1828, 2013.

[9] Attila Szabó, Qiyang Hu, Tiziano Portenier, Matthias Zwicker, and Paolo Favaro, "Understanding degeneracies and ambiguities in attribute transfer," in *Proceedings of ECCV*, 2018, vol. 11209, pp. 721–736.

[10] Diederik P. Kingma and Max Welling, "Auto-encoding variational bayes," in *Proceedings of ICLR*, 2014.

[11] Irina Higgins, Loic Matthey, et al., "beta-vae: Learning basic visual concepts with a constrained variational framework," in *Proceedings of ICLR*, 2017.

[12] Ricky T. Q. Chen, Xuechen Li, Roger Grosse, and David Duvenaud, "Isolating sources of disentanglement in variational autoencoders," in *Proceedings of ICLR*, 2018.

[13] Francesco Locatello, Stefan Bauer, et al., "Challenging common assumptions in the unsupervised learning of disentangled representations," in *Proceedings of ICLR*, 2019.

[14] Diane Bouchacourt, Ryota Tomioka, and Sebastian Nowozin, "Multi-level variational autoencoder: Learning disentangled representations from grouped observations," in *Proceedings of AAAI*, 2018, pp. 2095–2102.

[15] Yunhao Ge, Sami Abu-El-Haija, Gan Xin, and Laurent Itti, "Zero-shot synthesis with group-supervised learning," in *Proceedings of ICLR*, 2021.

[16] Xiaosong Wang, Yifan Peng, et al., "Chestx-ray8: Hospital-scale chest x-ray database and benchmarks on weakly-supervised classification and localization of common thorax diseases," in *Proceedings of CVPR*, 2017, pp. 3462–3471.

[17] Bo Zhou, Yuemeng Li, and Jiangcong Wang, "A weakly supervised adaptive densenet for classifying thoracic diseases and identifying abnormalities," *CoRR*, vol. abs/1807.01257, 2018.

[18] Pranav Rajpurkar, Jeremy Irvin, et al., "Chexnet: Radiologist-level pneumonia detection on chest x-rays with deep learning," *CoRR*, vol. abs/1711.05225, 2017.

[19] Ken C. L. Wong, Mehdi Moradi, Joy T. Wu, and Tanveer F. Syeda-Mahmood, "Identifying disease-free chest x-ray images with deep transfer learning," *Medical Imaging: Computer-Aided Diagnosis*, vol. 10950, pp. 109500P, 2019.

[20] Hongyu Wang and Yong Xia, "Chestnet: A deep neural network for classification of thoracic diseases on chest radiography," *CoRR*, vol. abs/1807.03058, 2018.

[21] Jeremy Irvin, Pranav Rajpurkar, et al., "Chexpert: A large chest radiograph dataset with uncertainty labels and expert comparison," in *Proceedings of AAAI*, 2019, pp. 590–597.

[22] Hieu H. Pham, Tung T. Le, Dat Q. Tran, Dat T. Ngo, and Ha Q. Nguyen, "Interpreting chest x-rays via cnns that exploit hierarchical disease dependencies and uncertainty labels," *Neurocomputing*, vol. 437, pp. 186–194, 2021.

[23] Li Yao, Eric Poblenz, Dmitry Dagunts, Ben Covington, Devon Bernard, and Kevin Lyman, "Learning to diagnose from scratch by exploiting dependencies among labels," *CoRR*, vol. abs/1710.10501, 2017.

[24] Wenwu Ye, Jin Yao, Hui Xue, and Yi Li, "Weakly supervised lesion localization with probabilistic-cam pooling," *CoRR*, vol. abs/2005.14480, 2020.